Cheminfomatics Assignment 2

William Juhl

December 2023

Contents

1   Introduction

The goal of this paper is to discuss the implementation of a program that produces chemically correct graph grammar rules for a given set of educt molecules and product molecules. We will first discuss the formal specification of what defines a chemically valid mapping, then we will discuss the implementation of said specification and finally we will discuss the results of the program.

2 Formal definition of a chemically valid mapping

A chemically valid mapping is defined as a mapping between a subset of the atoms of the educt molecules and a subset of the atoms of the product molecules. It must also also create a singular cycle of even length. This cycle is constructed by XORing the bonds on the educt side with the bonds on the product side while adheering to the following rules:

1. If a bond is present on the educt side but not on the product side, it must be a single bond. This is called a *deletion*.

2. If a bond is present on the product side but not on the educt side, it must also be a single bond. This is called an *addition*.

3. If a bond exists between two atoms on both sides the difference in the bonds must not differ by more than one. If the bond on the product side is greater than that on the educt side it is called an *addition* and if the bond on the product side is less than that on the educt side it is called a *deletion*.

4. For each atom in the cycle it must have exactly one bond that is an addition and one bond that is a deletion.

If all of the above criteria are met, the mapping is considered chemically valid. We can further controll the mapping we are looking for by introducing two parameters, k and c. Here c describes the desired length of the cycle and k describes the context of the mapping. The context is defined as the atoms on both sides which do not have changing bonds and with a distance no more than k from the cycle. In general k controls the level of how generic the mapping is, with smaller k making the mapping more generic and bigger k making the mapping more specific.

3   Implementation

In this section we will discuss the implementation of the program. First we will discuss how to generate the morphisms which map the atoms from the educt to the atoms on the product. Then we will discuss how we can use the morphism to construct a new graph by XORing the bonds on the educt side with the bonds on the product side. Then we will expand the morphism to include the desired context based on the parameter k. Finally we will discuss how we can detect, construct and validate a potential cycle which indicates a chemically valid mapping.

3.1   Morphism Generation

To lazily generate morphisms we use two generators, *Morphism_Generator* which employs a helper generator *Combinatorial_Generator*. The generator picks c atoms from the educt and for each unique atom pick the same amount of atoms from the product. This ensures that we have a one-to-one mapping between the atoms from the educt and the atoms from the product and also that we do not have a morphism which maps a Hydrogen atom to a Carbon atom.

3.1.1   Combinatorial_Generator

To pick n amount of atoms from a set of atoms we use a bit field where exactly n bits are set and a one indicates that the atom is picked. To permute the next combination of n atoms we simply permute the bit field. If no more unique combinations can be generated return false, otherwise return true.

3.1.2   Morhpism_Generator

By using the above generator we pick c atoms from the educt. For each of the unique atoms from the educt we pick exactly the same amount of atoms from the product. Then constructing a morphism is as simple as mapping the index of the atom from the educt to the index of the atom from the product. We can generate the next morphism by picking the next combination on the product side. If there are no more unique combinations on the product side we pick a new combination on the educt side and start over. Once both sides have no more unique combinations we have generated all possible morphisms.

3.1.3   Observations

From observing and trying different orderings for which atoms are picked initially, we can gain a huge performance boost by mapping Hydrogen atoms last. Because Hydrogen atoms only have a single bond they are not very likely to be part of a cycle. For that reason we can pick them last and thus reduce the number of morphisms we have to generate before finding a valid cycle. To do this we have opted to pick the atoms based on their weight making heavier atoms being picked first for the cycle.

3.2 Educt XOR Product

When we have a morphism we can construct a new graph by XORing the bonds on the educt side with the bonds on the product side. This is done by iterating through each atom pair, *left* and *right*, in the morphism and checking for the following:

1. If the left atom has a bond to another atom in the morphism and the right atom does **not** have a bond to another atom in the morphism we have the following two cases:

   (a) If it is a single bond we add an edge with label REMOVE between the left atom's source and target.
   (b) If the bond is a double or a triple bond it is an invalid edge.

2. If the left atom does **not** have a bond to another atom in the morphism and the right atom does have a bond to another atom in the morphism we have the following two cases:

   (a) If it is a single bond we add an edge with label ADD between the right atom's source and target.
   (b) If the bond is a double or a triple bond it is an invalid edge.

3. If both the left and the right have a bond to another atom in the morphism we have the following five cases:

   (a) If the bond on the left is a single bond and the bond on the right is a double add the edge with label ADD.
   (b) If the bond on the left is a double bond and the bond on the right is a single bond add the edge with label REMOVE.
   (c) if the bond on the left is a double bond and the bond on the right is a triple bond add the edge with label ADD.
   (d) if the bond on the left is a triple bond and t he bond on the right is a double bond add the edge with label REMOVE.
   (e) Otherwise it is an invalid edge

The resulting graph is the XORed graph of the educt and the product using the given morphism. If any of the edges are invalid, the morphism is also not chemically valid.

3.3 Cycle Detection and Construction

By doing a Depth-First-Search on the XORed graph we can detect cycles by registering when a back edge is found. When a back edge is found we can construct a cycle by following the path from the back edge to the current node.

3.4 Cycle validation

If the above step produces more than a single cycle or if no cycle were detected we discard this morphism. Otherwise we can validate the cycle by checking if the cycle is of even length and if the cycle contains exactly one addition and one deletion for each atom in the cycle.

3.5   Context Expansion

To add the context defined by the parameter k we simply expand the morphism to include all atoms within k bonds of the cycle which adhere to both of the two following conditions:

1. The bond must be the same on the educt side as on the product side

2. The atom must be connected to the morphism on both sides.

This ensures that we only expand the morphism to include atoms which can reach the cycle within k bonds and which share the same bond strength on the educt and the product side.

4 Results

We will now compute a series of tests to see how the program performs. I am running the program inside a Docker container with the Rosetta translation layer on my Macbook Pro running a 10 core M1 Pro processor with access to 8 GB of RAM out of a total of 16 GB of RAM.

Educts	Products	k	c	n	t	s
C=C, C=C	C1CCC1	0	4	2	10	0.250s
C=C, C=C	C1CCC1	1	4	2	10	0.255s
O, Cl, CC(=O)OCC	Cl, OCC, CC(=O)O	0	6	3	10	0.804s
O, Cl, CC(=O)OCC	Cl, OCC, CC(=O)O	1	6	3	10	0.804s
O, Cl, CC(=O)OCC	Cl, OCC, CC(=O)O	2	6	3	10	0.811s
C1C(O)CC(O)C(O)C1	C=CO, C=CO, C=CO	0	6	4	10	0.437s
C1C(O)CC(O)C(O)C1	C=CO, C=CO, C=CO	1	6	6	10	0.455s
CC=CC=CC, OC1C=CC=CC=1	C=CC=CC=C, OC(=C)C=CC=C	0	6	4	10	4.169s
CC=CC=CC, OC1C=CC=CC=1	C=CC=CC=C, OC(=C)C=CC=C	1	6	6	10	4.163s
CC=CC=CC, OC1C=CC=CC=1	C=CC=CC=C, OC(=C)C=CC=C	2	6	7	10	4.146s
CC, OC1C=CC=CC=1	C=C, OC(=C)C=CC=C	0	6	4	10	3.266s
CC, OC1C=CC=CC=1	C=C, OC(=C)C=CC=C	1	6	4	10	3.264s
CC, OC1C=CC=CC=1	C=C, OC(=C)C=CC=C	2	6	4	10	3.290s
OP(=O)(O)OP(=O)(O)O, O	O=P(O)(O)O, O=P(O)(O)O	0	4	1	10	0.322s
OP(=O)(O)OP(=O)(O)O, O	O=P(O)(O)O, O=P(O)(O)O	1	4	1	10	0.317s
OP(=O)(O)OP(=O)(O)O, O	O=P(O)(O)O, O=P(O)(O)O	0	6	1	10	0.399s
OP(=O)(O)OP(=O)(O)O, O	O=P(O)(O)O, O=P(O)(O)O	1	6	1	10	0.397s
C#N, C#N	N=CC#N	0	4	1	2	0.254s
C#N, C#N	N=CC#N	1	4	1	2	0.270s

Table 1: k is the context, c is the cycle length, n is the number of unique rules found and t is the number of valid morphisms generated. The last column is the *real* time it took to generate the rules. The time is measured in seconds.

5 Example

To illustrate how a difference in context will affect the derivations, we have constructed the following example.

From the above table, derive the rules with cycle length of 6 from the third row. We have however removed the Chlorid from the educt and the product side. Then by adding to the subset only the educts and applying the generated rules once will result in the below figure. We will use context arguments of 0 and 1 and depict them below.

As we can observe (a) contains a bigger hypergraph than that in (b).

Figure 1: Hypergraphs from running a derivation once with the generated rules.
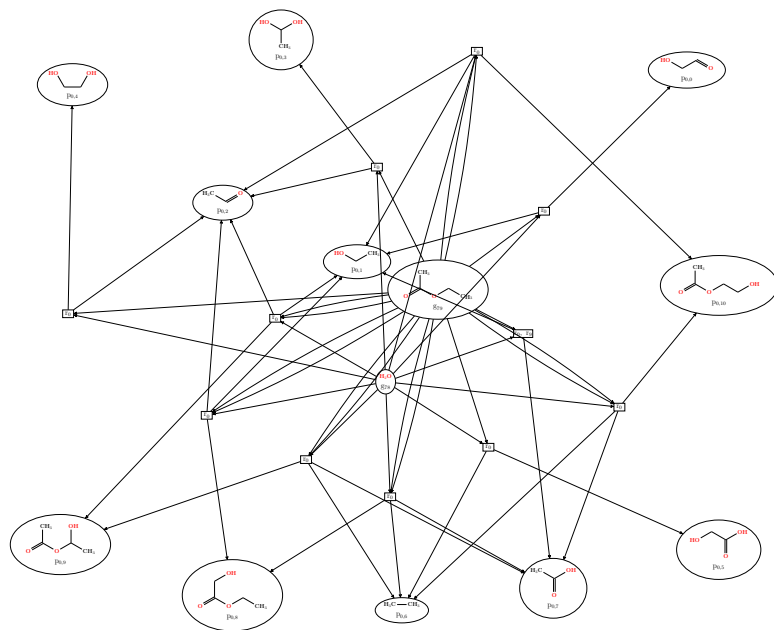
(a) Context k = 0
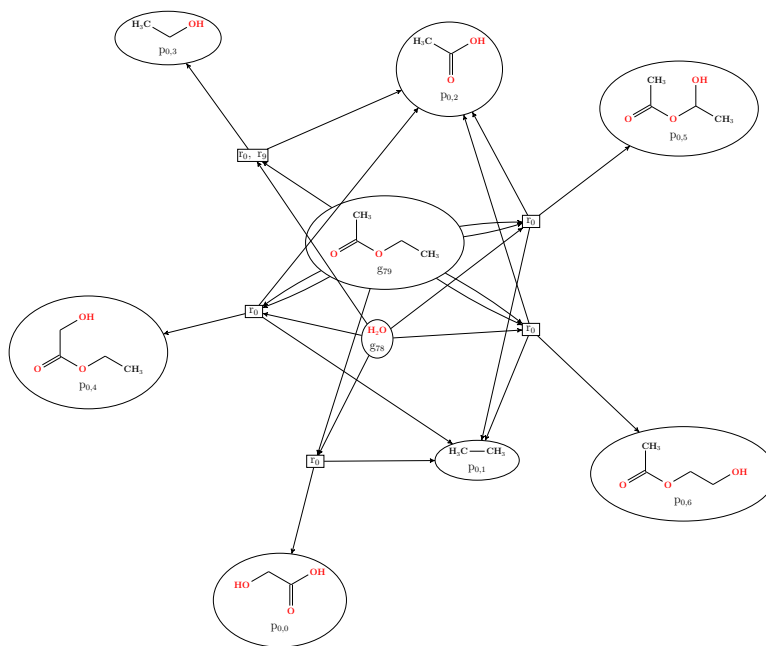
(b) Context k = 1

6 Improvements

The main improvement that can be done is to speed up finding valid morphisms resulting in unique rules. Currently we are generating all possible morphisms without checking for symmetry. My hypothesis is that many of the morphisms we generate result in the exact same rules. This can be seen by the number of discarded rules from the table above. By pruning the morphisms from the same cycle which result in an identical rule we can possibly get a more diverse set of rules.

If a cycle has been detected with even length and only a few edges being of the wrong type, we could potentially find another vertex which would satisfy the cycle. This would decrease the amount of morphisms we have to generate before finding a valid cycle.

Another way to generate morphisms would be to pick a starting atom on the educt and product side. Then as long as the desired cycle length has not yet been reached, add an edge to another vertex from the current mapping alternating between adding the edge from the educt and product. This should be used in conjunction with a backtracking algorithm to fix the issue of the chain of atoms not resulting in a cycle. By doing this we ensure at the same time that the cycle exists, is of even length and is alternating between addition and deletion. At the same time we can also possibly construct a more diverse set of rules by choosing different starting atoms for future morphisms.

As a final note our implementation has some steps which could be squashed into a single step. For example when we detect a cycle, we could at the same time do a validity check. This would save us from having to construct the cycle and then check if it is valid. And if the above suggestion of searching for a cycle while constructing the morphism, we could also remove XORing of edges step as it would be implemented in the morphism generation.